

BLOG



SQLAlchemy ORM Tutorial for Python Developers

Let's learn how to use SQLAlchemy ORM to persist and query data on Python applications.



Bruno Krebs

November 09, 2017

0 0 20

TL;DR: In this article, we will learn how to use SQLAlchemy as the ORM (Object Relational Database) library to communicate with relational database engines. First, we will learn about some core concepts of SQLAlchemy (like engines and connection pools), then we will learn how to map Python classes and its relationships to database tables, and finally we will learn how to retrieve (query) data from these tables. The code snippets used in this article can be found in this GitHub repository.

"Learn how to use SQLAlchemy ORM to persist and query data on Python applications."

TWEET THIS 

USE AUTH0 FOR F

SQLAlchemy is a library that facilitates the communication between Python programs and

Subscribe to more awesome content. databases. Most of the times, this library is used as an Object Relational Mapper (ORM) tool that

translates Python classes to tables on relational databases and automatically converts function

er provides a standard interface that allows developers to

create database-agnostic code to communicate with a wide variety of database engines.


Related Posts

As we will see in this article, SQLAlchemy relies on common design patterns (like Object Pools)

to now **Developing RESTful AP** developers to create and ship enterprise-grade, production-ready applications easily. **with Python and Flask**

Besides that, with SQLAlchemy, boilerplate code to handle tasks like database connections is

abstracted away to let developers focus on business logic.



vided by SQLAlchemy, we need to learn how the core
nce important concepts that every Python developer
SQLAlchemy applications.

(Base API) was created to specify how Python modules
se their interfaces. Although we won't interact with this
s a facade to it—it's good to know that it defines how
, `commit`, and `rollback` must behave. Consequently,

whenever we use a Python module that adheres to the specification, we can rest assured that we
will find these functions and that they will behave as expected.

In this article, we are going to install and use the most popular PostgreSQL DBAPI

implementation available: `psycopg`. Other Python drivers communicate with PostgreSQL as

well, but `psycopg` is the best candidate since it fully implements the DBAPI specification and
has great support from the community.

To better understand the DBAPI specification, what functions it requires, and how these

functions behave, take a look into the Python Enhancement Proposal that introduced it. Also, to

learn about what other database engines we can use (like MySQL or Oracle), take a look at the

SQLAlchemy Engines

Subscribe to more awesome conte

Whenever we want to use SQLAlchemy to interact with a database, we need to create an *Engine*.


Managing two crucial factors: *Pools* and *Dialects*. The

following two sections will explain what these two concepts are, but for now it suffices to say that

SQLAlchemy uses them to interact with DBAPI functions.

To create an engine and start interacting with databases, we have to import the `create_engine` function from the `sqlalchemy` library and issue a call to it:

Developing RESTful AP
with Python and Flask
Bruno Krebs



```
engine = create_engine('mysql://usr:pass@localhost:5432/sqlalchemy')
```

to communicate with an instance running locally on
es that it will use `usr` and `pass` as the credentials to
Note that, creating an engine does *not* connect to the
ned to when it's needed (like when we submit a query,

Since SQLAlchemy relies on the DBAPI specification to interact with databases, the most common database management systems available are supported. PostgreSQL, MySQL, Oracle, Microsoft SQL Server, and SQLite are all examples of engines that we can use alongside with SQLAlchemy. To learn more about the options available to create SQLAlchemy engines, take a look at the official documentation.

SQLAlchemy Connection Pools

Connection pooling is one of the most traditional implementations of the object pool pattern.

Object pools are used as caches of pre-initialized objects ready to use. That is, instead of spending time to create objects that are frequently needed (like connections to databases) the

USE AUTH0 FOR F

The main reason why programs take advantage of this design pattern is to improve performance.

Subscribe to more awesome content

In the case of database connections, opening and maintaining new ones is expensive,

time-consuming and wastes resources. Besides that, this pattern allows easier management of



ation might use simultaneously.

There are various implementations of the connection pool pattern available on SQLAlchemy .

Related Posts


For example, creating an Engine through the create_engine() function usually generates a

Q. Developing RESTful AP
- Python and Flask
pool size of 5 connections.



Bruno Krebs

As usual production-ready programs need to override these defaults (to fine-tune pools to their



tions of connection pools provide a similar set of

shows the most common options with their

ections that the pool will handle.

exceeding connections (relative to pool_size) the pool

um age (in seconds) of connections in the pool.

seconds the program will wait before giving up on

getting a connection from the pool.

To learn more about connection pools on SQLAlchemy, check out the official documentation.

SQLAlchemy Dialects

As SQLAlchemy is a facade that enables Python developers to create applications that communicate to different database engines through the same API, we need to make use of *Dialects*. Most of the popular relational databases available out there adhere to the SQL (Structured Query Language) standard, but they also introduce proprietary variations. These

USE AUTH0 FOR F


For example, let's say that we want to fetch the first ten rows of a table called `people`. If our data was being held by a Microsoft SQL Server database engine, SQLAlchemy would need to issue the following query:



Related Posts `SELECT TOP 10 * FROM people;`

Developing RESTful AP
 with Python and Flask
 Bruno Krebs

our data was persisted on MySQL instance, then SQLAlchemy would need to issue:



to issue, SQLAlchemy needs to be aware of the type of
 exactly what *Dialects* do. They make SQLAlchemy

following list of dialects:

Microsoft SQL Server

MySQL

Oracle

PostgreSQL

SQLite

Sybase

Dialects for other database engines, like Amazon Redshift, are supported as external projects but can be easily installed. Check out the official documentation on SQLAlchemy Dialects to learn more.

USE AUTH0 FOR F

ORM, which stands for *Object Relational Mapper*, is the specialization of the *Data Mapper*

Subscribe to more awesome conte

design pattern that addresses relational databases like MySQL, Oracle, and PostgreSQL. As

explained by Martin Fowler in the article, *Mappers* are responsible for moving data between

in independent of each other. As object-oriented

programming languages and relational databases structure data on different ways, we need

Related Posts

specific code to translate from one schema to the other.

Developing RESTful AP
For example, in a programming language like Python, we can create a `Product` class and an

with Python and Flask
class to relate as many instances as needed from one class to another (i.e. `Product` can

contain a list of instances of `Order` and vice-versa). Though, on relational databases, we need

products, another one to persist orders, and a third one to

QLAlchemy ORM is an excellent *Data Mapper* solution
es and to move data between instances of these classes

Learn More
sured that we will get support for the most common

For example, booleans, dates, times, strings, and

numeric values are a just a subset of the types that SQLAlchemy provides abstractions for.

Besides these basic types, SQLAlchemy includes support for a few vendor-specific types (like

JSON) and also allows developers to create custom types and redefine existing ones.

To understand how we use SQLAlchemy data types to map properties of Python classes into
columns on a relation database table, let's analyze the following example:

```
class Product(Base):
    __tablename__ = 'products'
    id=Column(Integer, primary_key=True)
    title=Column('title', String(32))
```

USE AUTH0 FOR F

```
price=Column('price', Numeric)
```

Subscribe to more awesome conte




defining a class called `Product` that has six properties. Let's take a look at what these properties do:

Related Posts

The `__tablename__` property tells SQLAlchemy that rows of the `products` table must be mapped to this class.

The `__primary__` property identifies that this is the `primary_key` in the table and that its type is

```
Integer
```



column in the table has the same name of the property

column in the table has the same name of the

column in the table has the same name of the

column in the table has the same name of the property

Seasoned developers will notice that (usually) relational databases do not have data types with these exact names. SQLAlchemy uses these types as generic representations to what databases support and use the dialect configured to understand what types they translate to. For example, on a PostgreSQL database, the `title` would be mapped to a `varchar` column.

SQLAlchemy Relationship Patterns

Now that we know what ORM is and have look into data types, let's learn how to use SQLAlchemy to map relationships between classes to relationships between tables. SQLAlchemy supports four types of relationships: One To Many, Many To One, One To One, and Many To Many.

USE AUTH0 FOR F

ORM in Practice action we will do a hands-on to practice mapping classes into

Subscribe to more awesome conte

tables and to learn how to insert, extract, and remove data from these tables.



The first type, *One To Many*, is used to mark that an instance of a class can be associated with

many instances of another class. For example, on a blog engine, an instance of the `Article` class

could be associated with many instances of the `Comment` class. In this case, we would map the


Defined classes and its relation as follows:



Developing RESTful AP
with Python and Flask

Bruno Krebs

```
class Article(Base):
```



```
=True)  
")
```

```
=True)  
eignKey('articles.id'))
```

the same relationship described above but from the

other perspective. To give a different example, let's say that we want to map the relationship

between instances of `Tire` to an instance of a `Car`. As many tires belong to one car and this

car contains many tires, we would map this relation as follows:


```
class Tire(Base):  
    __tablename__ = 'tires'  
    id = Column(Integer, primary_key=True)  
    car_id = Column(Integer, ForeignKey('cars.id'))  
    car = relationship("Car")
```




Related Posts



```
class Person(Base):
```



```

    =True)

    filePhone", uselist=False, back_populates="person"

    =True)

   ignKey('people.id'))

    back_populates="mobile_phone")

```

eters to the relationship function. The first one,

09/04/2018 20:18

Subscribe to more awesome conte

```
students_classes_association = Table('students_classes', Base.metadata,
```

```
    Column('student_id', Integer, ForeignKey('students.id')),
```



```
    Column('classes_id', Integer, ForeignKey('classes.id'))
```

Related Posts

[class Student\(Base\):](#)

```
    __tablename__ = 'students'
```

```
    id = Column(Integer, primary_key=True)
```




```
    classes = relationship("Class", secondary=students_classes_association)
```

Bruno Krebs

[class Class\(Base\):](#)

```
    __tablename__ = 'classes'
```

```
    id = Column(Integer, primary_key=True)
```



to persist the association between instances of `Student` and `Class` wouldn't be possible without an extra table. Note that, to create this table, we passed it in the `secondary` parameter of the `relationship` function.

Let's explore some of the mapping options supported by SQLAlchemy.

In the next section, we'll take a more in-depth look into each one of the

available relationship patterns. Besides that, [the official documentation is a great reference to learn more about relationship patterns on SQLAlchemy.](#)

SQLAlchemy ORM Cascade

Whenever rows in a particular table are updated or deleted, rows in other tables might need to suffer changes as well. These changes can be simple updates, which are called cascade updates, or full deletes, known as cascade deletes. For example, let's say that we have a table called

`shopping_carts`, a table called `products`, and a third one called `shopping_carts_products`

that connects the first two tables. If, for some reason, we need to delete rows from

`shopping_carts` we will need to delete the related rows from `shopping_carts_products` as well.

USE AUTH0 FOR F

To make this kind of operation easy to maintain, SQLAlchemy ORM enables developers to map

Subscribe to more awesome content
 cascade behavior when using `Relationship()` constructs. Like that, when operations are

performed on parent objects, child objects get updated/deleted as well. The following list

used cascade strategies on SQLAlchemy ORM:

Related Posts

save-update : Indicates that when a parent object is saved/updated, child objects are

saved/updated as well.



Developing RESTful AP
 with Python and Flask

delete : Indicates that when a parent object is deleted, children of this object will be deleted

Bruno Krebs

as well.

delete-orphan : Indicates that when a child object loses reference to a parent, it will get

Join the Auth0
 Ambassador Program



[Learn More](#)

relations propagate from parent to children.

needed, the [SQLAlchemy documentation](#) provides an

implementation of the [Unit of Work](#) design pattern. As

explained by Martin Fowler, a Unit of Work is used to maintain a list of objects affected by a business transaction and to coordinate the writing out of these changes. This means that all modifications tracked by Sessions (Units of Works) will be applied to the underlying database together, or none of them will. In other words, Sessions are used to guarantee the database consistency.

The official SQLAlchemy ORM documentation about [Sessions](#) gives a great explanation how changes are tracked, how to get sessions, and how to create ad-hoc sessions. However, in this article, we will use the most basic form of session creation:

USE AUTH0 FOR F

Subscribe to more awesome conte

create an engine

engine = create_engine('postgresql://usr:pass@localhost:5432/sqlalchemy')

create a configured "Session" class

Session = sessionmaker(bind=engine)

Related Posts

create a Session


session = Session()



Developing RESTful AP with Python and Flask

Bruno Krebs

As we can see from the code snippet above, we only need one step to get sessions. We need to



the SQLAlchemy engine. After that, we can just issue
ions.

of the most important pieces of SQLAlchemy, it's time to
as, we will create a small project based on pipenv —a
ome classes to it. Then we will map these classes to
and learn how to query data.

Starting the Tutorial Project

To create our tutorial project, we have to have Python installed on our machine and `pipenv` installed as a global Python package. The following commands will install `pipenv` and set up the project. These commands are dependent on Python, so be sure to have it installed before proceeding:

install pipenv globally

pip install pipenv

create a new directory for our project

USE AUTH0 FOR F

```
# change working directory to it
```

cd sqlalchemy-tutorial
Subscribe to more awesome conte



Related Posts

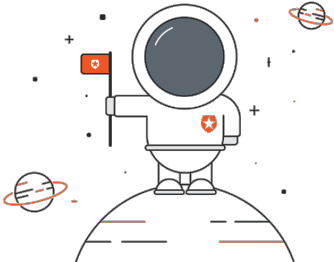
Running PostgreSQL

Developing RESTful AP

T with Python and Flask

Bruno Krebs
need a database to support our examples. As already mentioned, SQLAlchemy provides support for many different databases engines, but the instructions that follow will focus on PostgreSQL.

Join the Auth0
Ambassador Program



Learn More

PostgreSQL. One of them is to use some cloud (both of them have free tiers). Another possibility is to use a local environment. A third option is to run a PostgreSQL

because it has the performance of an instance and it's easy to create and destroy Docker instances. The first option is to install Docker locally.

create and destroy *dockerized* PostgreSQL instances with

the following commands:

```
# create a PostgreSQL instance
docker run --name sqlalchemy-orm-psql \
  -e POSTGRES_PASSWORD=pass \
  -e POSTGRES_USER=usr \
  -e POSTGRES_DB=sqlalchemy \
  -p 5432:5432 \
  -d postgres

# stop instance
docker stop sqlalchemy-orm-psql
```

Subscribe to more awesome conte

The first command, the one that creates the PostgreSQL instance, contains a few parameters that



Related Posts

`--name` : Defines the name of the Docker instance.

`-e POSTGRES_PASSWORD` : Defines the password to connect to PostgreSQL.



Developing RESTful AP


with Python and Flask

PostgreSQL user to connect to PostgreSQL.


Bruno Krebs

`-e POSTGRES_DB` : Defines the main (and only) database available in the PostgreSQL instance.

`-p 5432:5432` : Defines that the local `5432` port will tunnel connections to the same port in



Join the Auth0
Ambassador Program



[Learn More](#)

er instance will be created based on the official

dependencies

y two packages: `sqlalchemy` and `psycopg2` . The first and the second one, `psycopg2` , is the PostgreSQL communicate with the database. To install these

dependencies, we will use `pipenv` as shown:

```
# install sqlalchemy and psycopg2
pipenv install sqlalchemy psycopg2
```

This command will download both libraries and make them available in our Python virtual environment. Note that to run the scripts that we are going to create, we first need to spawn the virtual environment shell. That is, before executing `python somescript.py` , we need to execute `pipenv shell` . Otherwise, Python won't be able to find the installed dependencies, as they are just available in our new virtual environment.

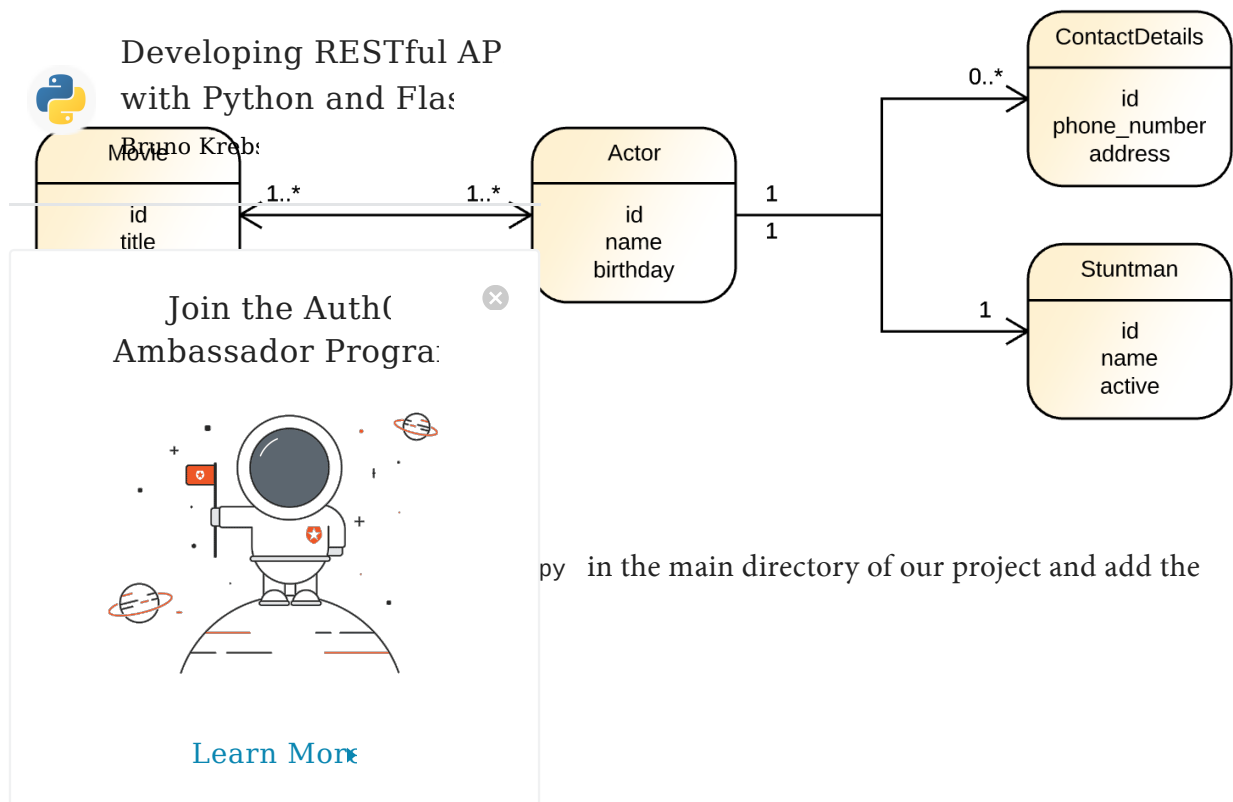
USE AUTH0 FOR F

After starting the *dockerized* PostgreSQL instance and installing the Python dependencies, we

can begin to map Python classes to database tables. In this tutorial, we will map four simple

Python classes to database tables: *Movie*, *Actor*, *ContactDetails*, and *Stuntman*. The following diagram

Related Posts



```
from sqlalchemy import create_engine
from sqlalchemy.ext.declarative import declarative_base
from sqlalchemy.orm import sessionmaker

engine = create_engine('postgresql://usr:pass@localhost:5432/sqlalchemy')
Session = sessionmaker(bind=engine)

Base = declarative_base()
```

This code creates:

and a base class for our classes definitions.

Subscribe to more awesome conte




ies. To do this, let's create a new file called `movie.py` and

add the following code to it:

Related Posts

coding=utf-8
 Developing RESTful AP
 with Python and Flask
 from sqlalchemy import Column, String, Integer, Date

from base import Base



=True)

e_date):

.date

The definition of this class and its mapping characteristics is quite simple. We start by making this class extend the `Base` class defined in the `base.py` module and then we add four properties to it:

1. A `__tablename__` to indicate what is the name of the table that will support this class.
2. An `id` to represent the primary key in the table.
3. A `title` of type `String`.
4. A `release_date` of type `Date`.

The next class that we will create and map is the `Actor` class. Let's create a file called `actor.py`

USE AUTH0 FOR F

Subscribe to more awesome conte
coding=utf-8



g, Integer, Date

```
from base import Base
```

Related Posts


Developing RESTful AP
with Python and Flask



Bruno Krebs

```
id = Column(Integer, primary_key=True)
```

```
name = Column(String)
```



y):

r to the previous one. The differences are that the
a birthday instead of a release_date , and that it
f movies .

As many movies can have many actors and vice-versa, we will need to create a *Many To Many* relationship between these two classes. Let's create this relationship by updating the `movie.py` file as follows:

```
# coding=utf-8
```

```
from sqlalchemy import Column, String, Integer, Date, Table, ForeignKey
from sqlalchemy.orm import relationship
```

```
from base import Base
```

```
movies_actors_association = Table(
```

USE AUTH0 FOR F

```
Column(actor_id, Integer, ForeignKey(actors.id))
```

Subscribe to more awesome conte



```
__tablename__ = 'movies'
```

Related Posts

```
id = Column(Integer, primary_key=True)
```

```
title = Column(String)
```

```
release_date = Column(Date)
```

```
actors = relationship("Actor", secondary=movies_actors_association)
```


```
Bruno Krebs
```

```
Bruno Krebs
```

```
def __init__(self, title, release_date):
```

```
self.title = title
```

Join the Auth0
Ambassador Program



[Learn More](#)

the previous one is that:

`relationship`, `ForeignKey`, and `relationship`;

an association table that connects rows of `actors` and rows of

`Movie` and configured the

`movies_actors_association` as the intermediary table.

The next class that we will create is `Stuntman`. In our tutorial, a particular `Actor` will have only one `Stuntman` and this `Stuntman` will work only with this `Actor`. This means that we need to create the `Stuntman` class and a *One To One* relationship between these classes. To accomplish that, let's create a file called `stuntman.py` and add the following code to it:

```
# coding=utf-8
```

```
from sqlalchemy import Column, String, Integer, Boolean, ForeignKey
```

USE AUTH0 FOR F

`from base import Base`

Subscribe to more awesome conte




Related Posts

`id = Column(Integer, primary_key=True)`
`name = Column(String)`

`active = Column(Boolean)`
`actor_id = Column(Integer, ForeignKey('actors.id'))`
`actor = relationship("Actor", backref=backref("stuntman", uselist=False))`

Developing RESTful AP
 with Python and Flask
 Bruno Krebs

`def __init__(self, name, active, actor):``self.name = name`

`actor` property references an instance of `Actor` and that
`stuntman` that is not a list (`uselist=False`). That is,
 an , SQLAlchemy will also load and populate the

up in our tutorial is `ContactDetails` . Instances of this
 class will hold a `phone_number` and an `address` of a particular `Actor` , and one `Actor` will be
 able to have many `ContactDetails` associated. Therefore, we will need to use the *Many To One*
 relationship pattern to map this association. To create this class and this association, let's create a
 file called `contact_details.py` and add the following source code to it:

`# coding=utf-8`

`from sqlalchemy import Column, String, Integer, ForeignKey`
`from sqlalchemy.orm import relationship`

`from base import Base`

USE AUTH0 FOR F

```
class ContactDetails(Base):
```

```
    tablename = 'contact_details'
```

Subscribe to more awesome conte


```
    id = Column(Integer, primary_key=True)
```

```
    address = Column(String)
```

```
    actor_id = Column(Integer, ForeignKey('actors.id'))
```


Related Posts

```
    actor = relationship("Actor", backref="contact_details")
```

 **Developing RESTful AP**
with Python and Flask

```
def __init__(self, phone_number, address, actor):
    self.phone_number = phone_number
    self.address = address
    self.actor = actor
```

Join the Auth0
Ambassador Program



[Learn More](#)

sociation is kinda similar to creating a *One To One* association. After we instructed SQLAlchemy not to use lists. This association to a single instance instead of a list of instances.

SQLAlchemy ORM

Let's create a file called `inserts.py` and generate some data in the database. In this file, let's add the following code:

```
# coding=utf-8

# 1 - imports
from datetime import date

from actor import Actor
from base import Session, engine, Base
from contact_details import ContactDetails
from movie import Movie
from stuntman import Stuntman


# 2 - generate database schema
```

USE AUTH0 FOR F

```
# 3 - Create a new session
```

session = Session()
Subscribe to more awesome conte

```
# 4 - create movies
```



B Identity", date(2002, 10, 11))

```
furious_7 = Movie("Furious 7", date(2015, 4, 2))
```

```
pain_and_gain = Movie("Pain & Gain", date(2013, 8, 23))
```

Related Posts

```
# 5 - creates actors
```

Developing RESTful AP
with Python and Flask
Bruno Krebs

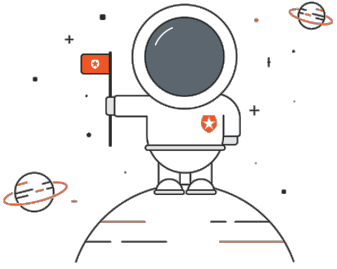
```
+t_damon = Actor("Matt Damon", date(1970, 10, 8))
```

```
wayne_johnson = Actor("Dwayne Johnson", date(1972, 5, 2))
```

```
mark_wahlberg = Actor("Mark Wahlberg", date(1971, 6, 5))
```

```
# 6 - add actors to movies
```

Join the Auth0
Ambassador Program



Learn More

×

```
on, mark_wahlberg]
```

```
55 2671", "Burbank, CA", matt_damon)
```

```
555 5623", "Glendale, CA", dwayne_johnson)
```

```
21 444 2323", "West Hollywood, CA", dwayne_johnso
```

```
33 9428", "Glendale, CA", mark_wahlberg)
```

```
True, matt_damon)
```

```
wayne_stuntman = Stuntman("John Roe", True, dwayne_johnson)
```

```
mark_stuntman = Stuntman("Richard Roe", True, mark_wahlberg)
```

```
# 9 - persists data
```

```
session.add(bourne_identity)
```

```
session.add(furious_7)
```

```
session.add(pain_and_gain)
```



```
session.add(matt_contact)
```

```
session.add(wayne_contact)
```

```
session.add(wayne_contact_2)
```

```
session.add(mark_contact)
```



```
session.add(matt_stuntman)
```

USE AUTH0 FOR F

10 - commit and close session

Subscribe to more awesome conte

`session.commit()``session.close()`

Related Posts split into 10 sections. Let's inspect them:

1. The first section implements the classes that we created, the SQLAlchemy engine, the Base class, the session factory, and date from the datetime module.




Developing RESTful AP

with Python and Flask

Bruno Krebs

2. The second section instructs SQLAlchemy to generate the database schema. This

generation occurs based on the declarations that we made while creating the four main



sion from the session factory.

ances of the `Movie` class.ces of the `Actor` class.ries. Note that the *Pain & Gain* movie references two *Vahlberg*.s of the `ContactDetails` class and defines what actors

ntmen and also defines what actors these stuntmen are

associated to.

9. The ninth section uses the current session to save the movies, actors, contact details, and stuntmen created. Note that we haven't explicitly saved actors. This is not needed because SQLAlchemy, by default, uses the save-update cascade strategy.
10. The tenth section commits the current session to the database and closes it.

To run this Python script, we can simply issue the `python inserts.py` command (let's not to run `pipenv shell` first) in the main directory of our database. Running it will create five tables in the PostgreSQL database and populate these tables with the data that we created. In the next section, we will learn how to query these tables.

USE AUTH0 FOR F

As we will see, querying data with SQLAlchemy ORM is quite simple. This library provides an intuitive, fluent API that enables developers to write queries that are easy to read and to

SQLAlchemy ORM, all queries start with a Query Object that is extracted from the `Session` object associated with a particular mapped class. To see this API in action,

let's create a file called `queries.py` and add to it the following source code:

Related Posts

Developing RESTful AP with Python and Flas




oding=utf-8

Bruno Kreb

1 - imports

from actor import Actor



```
print(f'{movie.title} was released on {movie.release_date}')
print('')
```

The code snippet above—that can be run with `python queries.py` ,—shows how easy it is to use SQLAlchemy ORM to query data. To retrieve all movies from the database, we just needed to fetch a session from the session factory, use it to get a query associated with `Movie` , and then call the `all()` function on this query object. The Query API provides dozens of useful functions like `all()` . In the following list, we can see a brief explanation about the most important ones:

`delete()` : Removes from the database the rows matched by a query.

`distinct()` : Applies a distinct statement to a query.

`exists()` : Returns a boolean value indicating if a subquery exists.


`first()` : Returns the first row in a query.

`related()` : Returns the row referenced by the primary key parameter passed as argument.

`join()` : Creates a SQL join in a query.

`limit()` : Limits the number of rows returned by a query.

`order_by()` : Sets an order in the rows returned by a query.



actions, let's append the following code to the

```
filter(Movie.release_date > date(2015, 1, 1)) \
.all()
```

```
print('### Recent movies:')
for movie in movies:
    print(f'{movie.title} was released after 2015')
print('')
```

```
# 6 - movies that Dwayne Johnson participated
the_rock_movies = session.query(Movie) \
    .join(Actor, Movie.actors) \
    .filter(Actor.name == 'Dwayne Johnson') \
    .all()
```


USE AUTH0 FOR F

```
print(f'The Rock starred in {movie.title}')
```

```
print('')
Subscribe to more awesome conte
```

```
# 7. Get actors that have houses in Glendale
y ) \


.join(ContactDetails) \

.filter(ContactDetails.address.ilike('%glendale%')) \
Related Posts
.all()
```

Developing RESTful AP
with Python and Flask

```
print('### Dwayne Johnson movies:')
actor in glendale_stars:
    Bruno Krebs
    print(f'{actor.name} has a house in Glendale')

print('')
```



as the `filter()` function to fetch only movies that
 . The sixth section shows how to use `join()` to fetch
 yne Johnson participated in. The seventh and last
`ilike()` functions to retrieve actors that have houses
 ython queries.py) now will result in the following

```
### All movies:
The Bourne Identity was released on 2002-10-11
Furious 7 was released on 2015-04-02
No Pain No Gain was released on 2013-08-23

### Recent movies:
Furious 7 was released after 2015

### Dwayne Johnson movies:
The Rock starred in No Pain No Gain
The Rock starred in Furious 7
```

USE AUTH0 FOR F

Mark Wariberg has a house in Glendale

Subscribe to more awesome conte




Forward and generates a code that is readable. To see other functions supported by the Query API, and their description, take a look at [the official documentation](#).

[Related Posts](#)

Developing RESTful APIs with Python and Flask
 "Querying data with SQLAlchemy ORM is easy and intuitive."
 Bruno Krebs



TWEET IT! KREBS



Auth0

easy and brings a lot of great features to the table. With a few lines of code to get:

, including [single sign-on](#)

(like Facebook, GitHub, Twitter, etc.)

(Active Directory, LDAP, SAML, etc.)

For example, to secure Python APIs written with Flask, we can simply create a `requires_auth` decorator:

```
# Format error response and append status code

def get_token_auth_header():
    """Obtains the access token from the Authorization Header"""
    auth = request.headers.get("Authorization", None)
    if not auth:
```

USE AUTH0 FOR F

`Authorization header is expected }, 401)`


Subscribe to more awesome conte

`parts = auth.split()``raise AuthError({"code": "invalid_header",``"description":`

Related Posts

`"Authorization header must start with"``" Bearer"}, 401)`Developing RESTful AP
with Python and Flask`raise AuthError({"code": "invalid_header",`

Bruno Krebs


`"description": "Token not found"}, 401)``elif len(parts) > 2:``raise AuthError({"code": "invalid_header",``"description": "Authorization header must be a valid token"}, 401)`Join the Auth0
Ambassador Program[Learn More](#)`token = get_token_auth_header()``jsonurl = urlopen("https://" + AUTH0_DOMAIN + "/.well-known/jwks.json")``jwks = json.loads(jsonurl.read())``unverified_header = jwt.get_unverified_header(token)``rsa_key = {}``for key in jwks["keys"]:``if key["kid"] == unverified_header["kid"]:``rsa_key = {``"kty": key["kty"],``"kid": key["kid"],``"use": key["use"],``"n": key["n"],``"e": key["e"]``}`

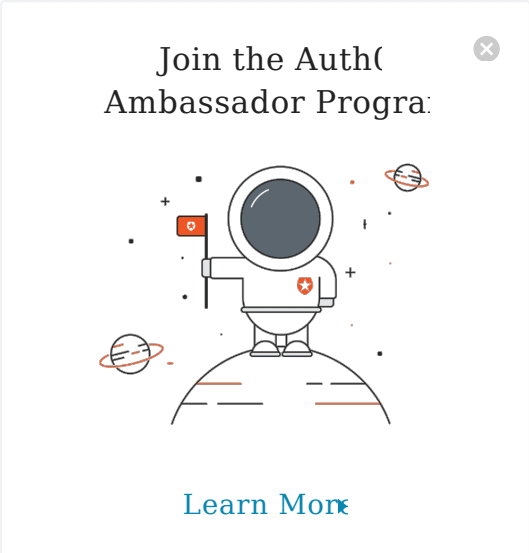
USE AUTH0 FOR F

payload = jwt.decode(
 token,
 Subscribe to more awesome conte
 rsa_key,
 algorithms=ALGORITHMS,
 audience=AUTH0_AUDIENCE,
 issuer="https://" + AUTH0_DOMAIN + "/"
)
 except jwt.ExpiredSignatureError:
 raise AuthError({"code": "token_expired",
 "description": "token is expired"}, 401)
 except jwt.JWTClaimsError:
 raise AuthError({"code": "invalid_claims",
 "description":
 "incorrect claims,"
 "please check the audience and issuer"}, 401)
 "code": "invalid_header",
 "description":
 "Unable to parse authentication"
 " token."}, 400)
 user = payload
)
 "valid_header",
 ": "Unable to find appropriate key"}, 400)

Related Posts

Developing RESTful AP
 with Python and Flask
 Bruno Krebs





Join the Auth0
 Ambassador Program

Learn More

Then use it in our endpoints:

```
# Controllers API

# This doesn't need authentication
@app.route("/ping")
@cross_origin(headers=['Content-Type', 'Authorization'])
def ping():
    return "All good. You don't need to be authenticated to call this"
```

USE AUTH0 FOR F

```
@cross_origin(headers=[ 'Content-Type', 'Authorization' ])
@requires_auth
def secured_ping():
    return "All good. You only get this message if you're authenticated"
```



Related Posts

For more about securing *Python APIs* with Auth0, take a look at this tutorial. Alongside with tutorials for backend technologies (like Python, Java, and PHP), the [Auth0 Docs](#) webpage **Developing RESTful APIs** provides tutorials for *Mobile/Native apps* and *Single-Page applications*.



Bruno Krebs

Next Steps

Join the Auth0 Ambassador Program

[Learn More](#)

article. We've learned about basic SQLAlchemy concepts. After that, we've learned about how SQLAlchemy patterns, Cascade strategies, and the Query API. In the exercise. In summary, we had the chance to learn and SQLAlchemy and SQLAlchemy ORM. In the next article, complement RESTful APIs with Flask—the Python

Authentication that just works.

USE AUTH0 FOR FREE

Any app. Any device. Hosted anywhere.

USE AUTH0 FOR F

Subscribe to more awesome conte
Join the discussion...



DISQUS (?)

Related Posts

[Lee Gaines](#) • 16 days ago

Lovely tutorial... Thank you so much! I am confused about something, though...
Developing RESTful AP



with Python and Flask

In this case, we had to create a helper table to persist the association between instances of Student and instances of Class, as this wouldn't be possible without an extra table.

.....

Join the Auth0
Ambassador Program

[Learn More](#)

in a many-to-many relationship?

es • 16 days ago

ra table, where would we keep the information about which
ich records on table B?

e a "customers" table and a "products" table. Different
oducts and, as such, there is a many-to-many association here.
no way to persist the information of which customer bought

ssociation, we could save on the "customers" table what
ould buy just one. So, it would be a many-to-one association
ducts".

Not sure if my explanation was clear enough, so I point you to this reference (I swear it was a coincidence that they use the same example :D).

^ | v • Reply • Share >

[Lee Gaines](#) ➔ Bruno S. Krebs • 16 days ago

That is very helpful... I'm starting to see the light... But still just a bit confused...

In a many-to-many relationship, wouldn't it still be possible to create a relationship between those two tables without bringing in a third? I imagine it'd be a huge mess with a whole bunch of columns, one for each transaction between a customer and a product for example.

Is the idea of JOINS to avoid that mess? Or is it actually impossible to create a many-to-many association without a new table?

Thank you for your patience :)

^ | v • Reply • Share >

[Bruno S. Krebs](#) Mod ➔ [Lee Gaines](#) • 16 days ago

USE AUTH0 FOR F

Subscribe to more awesome conte



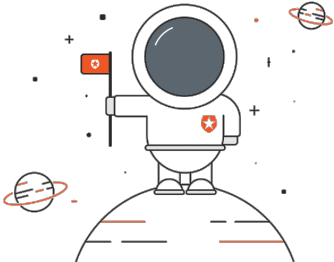
Related Posts



Developing RESTful AP with Python and Flas

Bruno Kreb

Join the Auth0
Ambassador Program



[Learn More](#)



PRODUCT

Pricing

Why Auth0

How It Works

COMPANY

About Us

Blog

USE AUTH0 FOR F

Press

Subscribe to more awesome conte

SECURITY

Security
Related Posts

White Hat



LE.

Developing RESTful AP
with Python and Flas

Bruno Krebs

Help & Support

Join the Auth0
Ambassador Program[Learn More](#)

CONTACT

10900 NE 8th St.

Suite 700

Bellevue, WA 98004

+1 (888) 235-2699

+1 (425) 312-6521

[Support Center](#)

Follow @auth0

12.7K followers



Mi piace 14.340

[Privacy Policy](#) [Terms of Service](#)

© 2013-2016 Auth0 © Inc. All Rights Reserved.